



THE LINUX FOUNDATION
OPEN SOURCE SUMMIT
EUROPE

Comparison of FOSS distributed filesystems

Marian "*HackMan*" Marinov, SiteGround

Who am I?

- Chief System Architect of SiteGround
- Sysadmin since 96
- Teaching Network Security & Linux System Administration in Sofia University
- Organizing biggest FOSS conference in Bulgaria
 - OpenFest

Why were we considering shared FS?

- We are a hosting provider
- We needed shared filesystem between VMs and containers
- Different size of directories and relatively small files
- Sometimes millions of files in a single dir

The story of our storage endeavors

- We started with DBRD + OCFS2
- I did a small test of cLVM + ATAoE
- We tried DRBD + OCFS2 for MySQL clusters
- We then switched to GlusterFS
- Later moved to CephFS
- and finally settled on good old NFS :)
but with the storage on Ceph RBD

Which filesystems were tested

- CephFS
- GlusterFS
- MooseFS
- OrangeFS
- BeeGFS - no stats

I haven't played with Lustre or AFS
And because of our history with OCFS2 and GFS2 I skipped them

Test cluster

- CPU i7-3770 @ 3.40GHz / 16G DDR3 1333MHz
- SSD SAMSUNG PM863
- 10Gbps Intel x520
- 40Gbps Qlogic QDR Infiniband
- IBM Blade 10Gbps switch
- Mellanox Infiniscale-IV switch

What did I test?

- Setup
- Fault tolerance
- Resource requirements (client & server)
- Capabilities (Redundancy, RDMA, caching)
- FS performance (creation, deletion, random read, random write)

Complexity of the setup

- CephFS - requires a lot of knowledge and time
- GlusterFS - relatively easy to setup and install
- OrangeFS - extremely easy to do basic setup
- MooseFS - extremely easy to do basic setup
- DRBD + NFS - very complex setup if you want HA
- BeeGFS -

CephFS

- Fault tollerant by desing...
 - until all MDSes start crashing
 - single directory may crash all MDSes
 - MDS behind on trimming
 - Client failing to respond to cache pressure
- Ceph has redundancy but lacks RDMA support

CephFS



- Uses a lot of memory on the MDS nodes
- Not suitable to run on the same machines as the compute nodes
- Small number of nodes 3-5 is a no go

CephFS tuning

- Placement Groups(PG)
- `mds_log_max_expiring` & `mds_log_max_segments` fixes the problem with trimming
- When you have a lot of inodes, increasing `mds_cache_size` works but increases the memory usage

CephFS fixes

- Client XXX failing to respond to cache pressure
 - This issue is due to the current working set size of the inodes.

```
ceph daemon mds.$(hostname) perf dump | grep inode
```

```
"inode_max": 100000, - max cache size value  
"inodes": 109488, - currently used
```

The fix for this is setting the `mds_cache_size` in `/etc/ceph/ceph.conf` accordingly.

CephFS fixes

- Client XXX failing to respond to cache pressure
 - Another “FIX” for the same issue is to login to current active MDS and run:

```
/etc/init.d/ceph restart mds
```

- This way it will change current active MDS to another server and will drop inode usage

GlusterFS

- Fault tolerance
 - in case of network hickups sometimes the mount may become inaccessible and requires manual remount
 - in case one storage node dies, you have to manually remount if you don't have local copy of the data
- Capabilities - local caching, RDMA and data Redundancy are supported. Offers different ways for data redundancy and sharding

GlusterFS

- High CPU usage for heavily used file systems
 - it required a copy of the data on all nodes
 - the FUSE driver has a limit of the number of small operations, that it can perform
- Unfortunately this limit was very easy to be reached by our customers

GlusterFS Tunning

- use distributed and replicated volumes instead of only replicated
 - gluster volume create VOL replica 2 stripe 2
- setup performance parameters

GlusterFS Tuning

volume set VOLNAME performance.cache-refresh-timeout 3

volume set VOLNAME performance.io-thread-count 8

volume set VOLNAME performance.cache-size 256MB

volume set VOLNAME performance.cache-max-file-size 300KB

volume set VOLNAME performance.cache-min-file-size 0B

volume set VOLNAME performance.readdir-ahead on

volume set VOLNAME performance.write-behind-window-size 100KB

GlusterFS Tuning

volume set VOLNAME features.lock-heal on
volume set VOLNAME cluster.self-heal-daemon enable
volume set VOLNAME cluster.metadata-self-heal on
volume set VOLNAME cluster.consistent-metadata on
volume set VOLNAME cluster.stripe-block-size 100KB
volume set VOLNAME nfs.disable on

FUSE tuning

FUSE

GlusterFS

entry_timeout

entry-timeout

negative_timeout

negative-timeout

attr_timeout

attribute-timeout

mount everything with “-o intr” :)

MooseFS



- Reliability with multiple masters
- Multiple metadata servers
- Multiple chunkservers
- Flexible replication (per directory)
- A lot of stats and a web interface

BeeGFS



- Metadata and Storage nodes are replicated by design
- FUSE based

OrangeFS

- No native redundancy uses corosync/pacemaker for HA
- Adding new storage servers requires stopping of the whole cluster
- It was very easy to break it

Ceph RBD + NFS

- the main tunnelling goes to NFS, by using the `cachefilesd`
- it is very important to have cache for both accessed and missing files
- enable FS-Cache by using `"-o fsc"` mount option

Ceph RBD + NFS

- verify that your mounts are with enabled cache:

```
# cat /proc/fs/nfsfs/volumes
```

NV	SERVER	PORT	DEV	FSID	FSC
v4	aaaaaaaa	801	0:35	17454aa0fddaa6a5:96d7706699eb981b	yes
v4	aaaaaaaa	801	0:374	7d581aa468faac13:92e653953087a8a4	yes

Sysctl tuning

`fs.nfs.idmap_cache_timeout = 600`

`fs.nfs.nfs_congestion_kb = 127360`

`fs.nfs.nfs_mountpoint_timeout = 500`

`fs.nfs.nlm_grace_period = 10`

`fs.nfs.nlm_timeout = 10`

Sysctl tuning

Tune the network card polling

```
net.core.netdev_budget=900
```

```
net.core.netdev_budget_usecs=1000
```

```
net.core.netdev_max_backlog=300000
```

Sysctl tuning

Increase network stack memory

```
net.core.rmem_max=16777216
```

```
net.core.wmem_max=16777216
```

```
net.core.rmem_default=16777216
```

```
net.core.wmem_default=16777216
```

```
net.core.optmem_max=16777216
```

Sysctl tuning

memory allocation min/pressure/max.

read buffer, write buffer, and buffer space

net.ipv4.tcp_rmem = 4096 87380 134217728

net.ipv4.tcp_wmem = 4096 65536 134217728

Sysctl tuning

turn off selective ACK and timestamps

```
net.ipv4.tcp_sack = 0
```

```
net.ipv4.tcp_timestamps = 0
```

```
net.ipv4.tcp_low_latency = 1
```

scalable or bbr

```
net.ipv4.tcp_congestion_control = scalable
```

Sysctl tuning

Increase system IP port range to allow for more concurrent connections

```
net.ipv4.ip_local_port_range = 1024 65000
```

OS tuning

```
vm.swappiness = 0
```

Increase system file descriptor limit

```
fs.file-max = 65535
```

Stats



- For small files, network BW is not a problem
- However network latency is a killer :(

Stats

Creation of 10000 empty files:

Local SSD	7.030s
MooseFS	12.451s
NFS + Ceph RBD	16.947s
OrangeFS	40.574s
Gluster distributed	1m48.904s

Stats

MTU	Congestion	result
MooseFS 10G		
1500	Scalable	10.411s
9000	Scalable	10.475s
9000	BBR	10.574s
1500	BBR	10.710s
GlusterFS 10G		
1500	BBR	48.143s
1500	Scalable	48.292s
9000	BBR	48.747s
9000	Scalable	48.865s

Stats

MTU	Congestion	result
MooseFS IPoIB		
1500	BBR	9.484s
1500	Scalable	9.675s
GlusterFS IPoIB		
9000	BBR	40.598s
1500	BBR	40.784s
1500	Scalable	41.448s
9000	Scalable	41.803s
GlusterFS RDMA		
1500	Scalable	31.731s
1500	BBR	31.768s

Stats

Creation of
10000
random size
files:

MTU	Congestion	result
MooseFS 10G		
1500	Scalable	3m46.501s
1500	BBR	3m47.066s
9000	Scalable	3m48.129s
9000	BBR	3m58.068s
GlusterFS 10G		
1500	BBR	7m56.144s
1500	Scalable	7m57.663s
9000	BBR	7m56.607s
9000	Scalable	7m53.828s

Stats

Creation of 10000 random size files:

MTU	Congestion	result
MooseFS IPoIB		
1500	BBR	3m48.254s
1500	Scalable	3m49.467s
GlusterFS RDMA		
1500	Scalable	8m52.168s



Thank you!

Marian *HackMan* Marinov
mm@siteground.com



Thank you! Questions?

Marian *HackMan* Marinov
mm@siteground.com



OPEN
EUROPE

SOURCE SUMMIT

THE LINUX FOUNDATION

